

Zaheer Quraishi

(+1) 514-953-7474

zaheerq-portfolio.netlify.app | zaheerqur1811@gmail.com | linkedin.com/in/zaheer-quraishi | github.com/zaheerqur

SUMMARY

Full-stack and data engineer with production experience across .NET, Java, React, and Azure, and end-to-end data pipelines spanning ingestion, transformation, orchestration, and ML model serving. **Open to relocation across Canada.**

EDUCATION

Concordia University
B.Eng. in Software Engineering

Montreal, QC, Canada
September 2021 – October 2025

EXPERIENCE

IT Developer

MHI R/J Aviation

May 2025 – Present

Boisbriand, QC, Canada

- Eliminated manual repository monitoring for engineering leadership by building a centralized **React 19 / TypeScript** dashboard aggregating commit activity and team membership context across **20+ Azure DevOps** repositories, with automated tech stack inference per repo, secured via **Azure AD / MSAL** with role-based access control.
- Replaced a manual, untracked supplier assessment process with a **Transport Canada-compliant C# / .NET** compliance module covering **9** request type workflows, automated document generation, status lifecycle tracking, and role-based access control, giving **100+** quality inspectors an auditable approval system for **500+** global aerospace suppliers.
- Restored data accessibility for newly onboarded client organizations by diagnosing and resolving a non-deterministic indexing failure in an **Azure Cognitive Search** and **SQL** indexing service managing **1.7 million** records.

Full-Stack Developer

BG Communications International

September 2024 – April 2025

Montreal, QC, Canada

- Continued development of a full-stack ERP web application in **Java / Spring Boot** and **Next.js** replacing a legacy desktop platform, taking ownership of the Resources module end-to-end across backend API design and frontend implementation.
- Designed and implemented data migration scripts and automated consistency checkers to transfer historical payroll and financial records from the legacy database to a normalized relational schema, verifying integrity through programmatic row count and field-level comparisons.
- Covered critical business logic with **Selenium** end-to-end test suites and **Mockito** unit tests, contributing to a **CI/CD** pipeline that ran automated checks on every pull request across a 10-person team.
- Led weekly stakeholder meetings to showcase progress and gather requirements, serving as the communication bridge between the development team and project supervisor throughout the engagement.

PROJECTS

NYC Taxi Analytics Platform | *Kafka, DuckDB, dbt, Airflow, FastAPI, scikit-learn, React, GitHub Actions* | Personal Project

- Built an **Airflow**-orchestrated event-driven data pipeline to ingest **9M+** NYC TLC trip records through **Kafka** into **DuckDB**, using a batched Python consumer to ensure **at-least-once delivery**.
- Transformed raw ingestion data into borough-level fare summaries and hourly demand heatmaps using **dbt**, with custom data quality tests enforced on every push via a **GitHub Actions** CI pipeline seeded with programmatic test data.
- Trained and deployed a **scikit-learn** fare prediction model served via a **FastAPI** endpoint, with a **React** dashboard visualizing live trip volume, revenue trends, and a real-time fare predictor.

EPL Match Outcome Predictor | *Python, scikit-learn, FastAPI, Docker, GitHub Actions, Render, Netlify* | Personal Project

- Built a self-sustaining **MLOps** pipeline that ingests **4,180** real Premier League matches across **11 seasons**, engineers **30** rolling features, trains and calibrates a **HistGradientBoostingClassifier**, and serves live predictions via a **FastAPI** REST API, designed to run indefinitely without manual intervention.
- Achieved **45.3%** accuracy on a fully held-out **2025–26** validation season with calibrated probability outputs, beating the home-win baseline and exposing real predictive signal from form, **xG** differentials, and head-to-head history.
- Automated the full retraining and deployment cycle via **GitHub Actions**: pulls fresh fixtures twice a week, retrains the model, pushes updated binaries to Git, and triggers a zero-downtime **Render** redeploy, ensuring the system updates itself every matchweek.

SKILLS

Languages: C#, SQL, TypeScript, JavaScript, Python, Java, HTML/CSS

Frameworks: ASP.NET MVC, .NET, ADO.NET, React.js, Spring Boot, Next.js, FastAPI, scikit-learn, pandas, dbt, Selenium, Tailwind CSS

Cloud & Platforms: Azure (Blob Storage, Key Vault, Active Directory, Cognitive Search, DevOps), Docker, Kafka, DuckDB, Airflow, Render, Netlify, Firebase, MSAL

Developer Tools: SSMS, Azure Storage Explorer, Git, GitHub Actions, JIRA, Confluence, Figma, VS Code, Visual Studio Enterprise, IntelliJ IDEA, PyCharm

Testing: Jest, React Testing Library, JUnit, Mockito, Selenium (Python), end-to-end test automation

Methodologies: Agile, Scrum, CI/CD, MLOps, ETL/ELT